

```
In [66]: # Monter Google Drive pour accéder à tes fichiers
from google.colab import drive
drive.mount('/content/drive')

# Copier Les fichiers depuis Google Drive vers L'environnement Colab
!cp -r /content/drive/MyDrive/nlp_lab/project1* .

# Installer Les dépendances depuis Le fichier requirements.txt
!pip install -r /content/drive/MyDrive/nlp_lab/project1/requirements.txt

# Installer Otter Grader (si nécessaire)
!pip install otter-grader

!pip install wget

# Fonction pour exécuter Les commandes shell
import os

def shell(commands, warn=True):
    """Exécute des commandes shell et affiche les résultats."""
    file = os.popen(commands)
    print(file.read().rstrip('\n'))
    exit_status = file.close()
    if warn and exit_status is not None:
        print(f"Command failed with exit code {exit_status}")
    return exit_status

# Vérifier si requirements.txt existe et télécharger le dépôt si nécessaire
shell("""
ls requirements.txt >/dev/null 2>&1
if [ ! $? = 0 ]; then
    rm -rf .tmp
    git clone https://github.com/cs236299-2024-winter/lab1-4.git .tmp
    mv .tmp/tests ./tests
    mv .tmp/requirements.txt ./requirements.txt
    rm -rf .tmp
fi
pip install -q -r requirements.txt
""")
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

Requirement already satisfied: torch in /usr/local/lib/python3.10/dist-packages (from -r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 1)) (2.5.1+cu121)

Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from -r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 2)) (3.8.0)

Requirement already satisfied: otter-grader==1.0.0 in /usr/local/lib/python3.10/dist-packages (from -r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (1.0.0)

Requirement already satisfied: transformers in /usr/local/lib/python3.10/dist-packages (from -r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 4)) (4.47.1)

Requirement already satisfied: datasets in /usr/local/lib/python3.10/dist-packages (from -r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 5)) (3.2.0)

Requirement already satisfied: tokenizers in /usr/local/lib/python3.10/dist-packages (from -r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 6)) (0.21.0)

Requirement already satisfied: pyyaml in /usr/local/lib/python3.10/dist-packages (from otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (6.0.2)

Requirement already satisfied: nbformat in /usr/local/lib/python3.10/dist-packages (from otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (5.10.4)

Requirement already satisfied: ipython in /usr/local/lib/python3.10/dist-packages (from otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (7.34.0)

Requirement already satisfied: nbconvert in /usr/local/lib/python3.10/dist-packages (from otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (7.16.4)

Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (4.67.1)

Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (75.1.0)

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (2.2.2)

Requirement already satisfied: tornado in /usr/local/lib/python3.10/dist-packages (from otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (6.3.3)

Requirement already satisfied: docker in /usr/local/lib/python3.10/dist-packages (from otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (7.1.0)

Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (3.1.4)

Requirement already satisfied: dill in /usr/local/lib/python3.10/dist-packages (from otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (0.3.8)

Requirement already satisfied: pdfkit in /usr/local/lib/python3.10/dist-packages (from otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (1.0.0)

Requirement already satisfied: PyPDF2 in /usr/local/lib/python3.10/dist-packages (from otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (3.0.1)

Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages

```
s (from torch->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 1)) (3.16.1)
Requirement already satisfied: typing-extensions>=4.8.0 in /usr/local/lib/python3.10/dist-packages (from torch->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 1)) (4.12.2)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 1)) (3.4.2)
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from torch->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 1)) (2024.9.0)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.10/dist-packages (from torch->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 1)) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from sympy==1.13.1->torch->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 1)) (1.3.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 2)) (1.3.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 2)) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 2)) (4.55.3)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 2)) (1.4.7)
Requirement already satisfied: numpy<2,>=1.21 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 2)) (1.26.4)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 2)) (24.2)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 2)) (11.0.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 2)) (3.2.0)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 2)) (2.8.2)
Requirement already satisfied: huggingface-hub<1.0,>=0.24.0 in /usr/local/lib/python3.10/dist-packages (from transformers->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 4)) (0.27.0)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.10/dist-packages (from transformers->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 4)) (2024.11.6)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from transformers->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 4)) (2.32.3)
Requirement already satisfied: safetensors>=0.4.1 in /usr/local/lib/python3.10/dist-packages (from transformers->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 4)) (0.4.5)
Requirement already satisfied: pyarrow>=15.0.0 in /usr/local/lib/python3.10/dist-packages (from datasets->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 5)) (17.0.0)
Requirement already satisfied: xxhash in /usr/local/lib/python3.10/dist-packages
```

```
(from datasets->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line
5)) (3.5.0)
Requirement already satisfied: multiprocessing<0.70.17 in /usr/local/lib/python3.10/
dist-packages (from datasets->-r /content/drive/MyDrive/nlp_lab/project1/requiem
ents.txt (line 5)) (0.70.16)
Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packages
(from datasets->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line
5)) (3.11.10)
Requirement already satisfied: aiohappyeyeballs>=2.3.0 in /usr/local/lib/python3.
10/dist-packages (from aiohttp->datasets->-r /content/drive/MyDrive/nlp_lab/proje
ct1/requirements.txt (line 5)) (2.4.4)
Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist
-packages (from aiohttp->datasets->-r /content/drive/MyDrive/nlp_lab/project1/req
uirements.txt (line 5)) (1.3.2)
Requirement already satisfied: async-timeout<6.0,>=4.0 in /usr/local/lib/python3.
10/dist-packages (from aiohttp->datasets->-r /content/drive/MyDrive/nlp_lab/proje
ct1/requirements.txt (line 5)) (4.0.3)
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-pa
ckages (from aiohttp->datasets->-r /content/drive/MyDrive/nlp_lab/project1/requir
ements.txt (line 5)) (24.3.0)
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dis
t-packages (from aiohttp->datasets->-r /content/drive/MyDrive/nlp_lab/project1/re
quirements.txt (line 5)) (1.5.0)
Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/d
ist-packages (from aiohttp->datasets->-r /content/drive/MyDrive/nlp_lab/project1/
requirements.txt (line 5)) (6.1.0)
Requirement already satisfied: propcache>=0.2.0 in /usr/local/lib/python3.10/dist
-packages (from aiohttp->datasets->-r /content/drive/MyDrive/nlp_lab/project1/req
uirements.txt (line 5)) (0.2.1)
Requirement already satisfied: yarl<2.0,>=1.17.0 in /usr/local/lib/python3.10/dis
t-packages (from aiohttp->datasets->-r /content/drive/MyDrive/nlp_lab/project1/re
quirements.txt (line 5)) (1.18.3)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-package
s (from python-dateutil>=2.7->matplotlib->-r /content/drive/MyDrive/nlp_lab/proje
ct1/requirements.txt (line 2)) (1.17.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python
3.10/dist-packages (from requests->transformers->-r /content/drive/MyDrive/nlp_la
b/project1/requirements.txt (line 4)) (3.4.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-pac
kages (from requests->transformers->-r /content/drive/MyDrive/nlp_lab/project1/re
quirements.txt (line 4)) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/di
st-packages (from requests->transformers->-r /content/drive/MyDrive/nlp_lab/proje
ct1/requirements.txt (line 4)) (2.2.3)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/di
st-packages (from requests->transformers->-r /content/drive/MyDrive/nlp_lab/proje
ct1/requirements.txt (line 4)) (2024.12.14)
Requirement already satisfied: jedi>=0.16 in /usr/local/lib/python3.10/dist-packa
ges (from ipython->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project
1/requirements.txt (line 3)) (0.19.2)
Requirement already satisfied: decorator in /usr/local/lib/python3.10/dist-packag
es (from ipython->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project
1/requirements.txt (line 3)) (4.4.2)
Requirement already satisfied: pickleshare in /usr/local/lib/python3.10/dist-pack
ages (from ipython->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/projec
t1/requirements.txt (line 3)) (0.7.5)
Requirement already satisfied: traitlets>=4.2 in /usr/local/lib/python3.10/dist-p
ackages (from ipython->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/pro
ject1/requirements.txt (line 3)) (5.7.1)
Requirement already satisfied: prompt-toolkit!=3.0.0,!<3.0.1,<3.1.0,>=2.0.0 in /u
```

```

usr/local/lib/python3.10/dist-packages (from ipython->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (3.0.48)
Requirement already satisfied: pygments in /usr/local/lib/python3.10/dist-packages (from ipython->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (2.18.0)
Requirement already satisfied: backcall in /usr/local/lib/python3.10/dist-packages (from ipython->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (0.2.0)
Requirement already satisfied: matplotlib-inline in /usr/local/lib/python3.10/dist-packages (from ipython->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (0.1.7)
Requirement already satisfied: pexpect>4.3 in /usr/local/lib/python3.10/dist-packages (from ipython->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (4.9.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (3.0.2)
Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (4.12.3)
Requirement already satisfied: bleach!=5.0.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (6.2.0)
Requirement already satisfied: defusedxml in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (0.7.1)
Requirement already satisfied: jupyter-core>=4.7 in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (5.7.2)
Requirement already satisfied: jupyterlab-pygments in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (0.3.0)
Requirement already satisfied: mistune<4,>=2.0.3 in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (3.0.2)
Requirement already satisfied: nbclient>=0.5.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (0.10.1)
Requirement already satisfied: pandocfilters>=1.4.1 in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (1.5.1)
Requirement already satisfied: tinycss2 in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (1.4.0)
Requirement already satisfied: fastjsonschema>=2.15 in /usr/local/lib/python3.10/dist-packages (from nbformat->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (2.21.1)
Requirement already satisfied: jsonschema>=2.6 in /usr/local/lib/python3.10/dist-packages (from nbformat->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (4.23.0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (2024.2)
Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from bleach!=5.0.0->nbconvert->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (0.5.1)
Requirement already satisfied: parso<0.9.0,>=0.8.4 in /usr/local/lib/python3.10/d

```

```

ist-packages (from jedi>=0.16->ipython->otter-grader==1.0.0->-r /content/drive/My
Drive/nlp_lab/project1/requirements.txt (line 3)) (0.8.4)
Requirement already satisfied: jsonschema-specifications>=2023.03.6 in /usr/local/lib/python3.10/dist-packages (from jsonschema>=2.6->nbformat->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (2024.10.1)
Requirement already satisfied: referencing>=0.28.4 in /usr/local/lib/python3.10/dist-packages (from jsonschema>=2.6->nbformat->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (0.35.1)
Requirement already satisfied: rpds-py>=0.7.1 in /usr/local/lib/python3.10/dist-packages (from jsonschema>=2.6->nbformat->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (0.22.3)
Requirement already satisfied: platformdirs>=2.5 in /usr/local/lib/python3.10/dist-packages (from jupyter-core>=4.7->nbconvert->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (4.3.6)
Requirement already satisfied: jupyter-client>=6.1.12 in /usr/local/lib/python3.10/dist-packages (from nbclient>=0.5.0->nbconvert->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (6.1.12)
Requirement already satisfied: ptyprocess>=0.5 in /usr/local/lib/python3.10/dist-packages (from pexpect>4.3->ipython->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (0.7.0)
Requirement already satisfied: wcwidth in /usr/local/lib/python3.10/dist-packages (from prompt-toolkit!=3.0.0,!<3.0.1,<3.1.0,>=2.0.0->ipython->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (0.2.13)
Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.10/dist-packages (from beautifulsoup4->nbconvert->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (2.6)
Requirement already satisfied: pyzmq>=13 in /usr/local/lib/python3.10/dist-packages (from jupyter-client>=6.1.12->nbclient>=0.5.0->nbconvert->otter-grader==1.0.0->-r /content/drive/MyDrive/nlp_lab/project1/requirements.txt (line 3)) (24.0.1)
Requirement already satisfied: otter-grader in /usr/local/lib/python3.10/dist-packages (1.0.0)
Requirement already satisfied: pyyaml in /usr/local/lib/python3.10/dist-packages (from otter-grader) (6.0.2)
Requirement already satisfied: nbformat in /usr/local/lib/python3.10/dist-packages (from otter-grader) (5.10.4)
Requirement already satisfied: ipython in /usr/local/lib/python3.10/dist-packages (from otter-grader) (7.34.0)
Requirement already satisfied: nbconvert in /usr/local/lib/python3.10/dist-packages (from otter-grader) (7.16.4)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from otter-grader) (4.67.1)
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from otter-grader) (75.1.0)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from otter-grader) (2.2.2)
Requirement already satisfied: tornado in /usr/local/lib/python3.10/dist-packages (from otter-grader) (6.3.3)
Requirement already satisfied: docker in /usr/local/lib/python3.10/dist-packages (from otter-grader) (7.1.0)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from otter-grader) (3.1.4)
Requirement already satisfied: dill in /usr/local/lib/python3.10/dist-packages (from otter-grader) (0.3.8)
Requirement already satisfied: pdfrkit in /usr/local/lib/python3.10/dist-packages (from otter-grader) (1.0.0)
Requirement already satisfied: PyPDF2 in /usr/local/lib/python3.10/dist-packages (from otter-grader) (3.0.1)
Requirement already satisfied: requests>=2.26.0 in /usr/local/lib/python3.10/dist-packages (from docker->otter-grader) (2.32.3)

```

Requirement already satisfied: urllib3>=1.26.0 in /usr/local/lib/python3.10/dist-packages (from docker->otter-grader) (2.2.3)

Requirement already satisfied: jedi>=0.16 in /usr/local/lib/python3.10/dist-packages (from ipython->otter-grader) (0.19.2)

Requirement already satisfied: decorator in /usr/local/lib/python3.10/dist-packages (from ipython->otter-grader) (4.4.2)

Requirement already satisfied: pickleshare in /usr/local/lib/python3.10/dist-packages (from ipython->otter-grader) (0.7.5)

Requirement already satisfied: traitlets>=4.2 in /usr/local/lib/python3.10/dist-packages (from ipython->otter-grader) (5.7.1)

Requirement already satisfied: prompt-toolkit!=3.0.0,!<3.0.1,<3.1.0,>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from ipython->otter-grader) (3.0.48)

Requirement already satisfied: pygments in /usr/local/lib/python3.10/dist-packages (from ipython->otter-grader) (2.18.0)

Requirement already satisfied: backcall in /usr/local/lib/python3.10/dist-packages (from ipython->otter-grader) (0.2.0)

Requirement already satisfied: matplotlib-inline in /usr/local/lib/python3.10/dist-packages (from ipython->otter-grader) (0.1.7)

Requirement already satisfied: pexpect>4.3 in /usr/local/lib/python3.10/dist-packages (from ipython->otter-grader) (4.9.0)

Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->otter-grader) (3.0.2)

Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader) (4.12.3)

Requirement already satisfied: bleach!=5.0.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader) (6.2.0)

Requirement already satisfied: defusedxml in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader) (0.7.1)

Requirement already satisfied: jupyter-core>=4.7 in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader) (5.7.2)

Requirement already satisfied: jupyterlab-pygments in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader) (0.3.0)

Requirement already satisfied: mistune<4,>=2.0.3 in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader) (3.0.2)

Requirement already satisfied: nbclient>=0.5.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader) (0.10.1)

Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader) (24.2)

Requirement already satisfied: pandocfilters>=1.4.1 in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader) (1.5.1)

Requirement already satisfied: tinycss2 in /usr/local/lib/python3.10/dist-packages (from nbconvert->otter-grader) (1.4.0)

Requirement already satisfied: fastjsonschema>=2.15 in /usr/local/lib/python3.10/dist-packages (from nbformat->otter-grader) (2.21.1)

Requirement already satisfied: jsonschema>=2.6 in /usr/local/lib/python3.10/dist-packages (from nbformat->otter-grader) (4.23.0)

Requirement already satisfied: numpy>=1.22.4 in /usr/local/lib/python3.10/dist-packages (from pandas->otter-grader) (1.26.4)

Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas->otter-grader) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->otter-grader) (2024.2)

Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas->otter-grader) (2024.2)

Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from bleach!=5.0.0->nbconvert->otter-grader) (0.5.1)

Requirement already satisfied: parso<0.9.0,>=0.8.4 in /usr/local/lib/python3.10/dist-packages (from jedi>=0.16->ipython->otter-grader) (0.8.4)

Requirement already satisfied: attrs>=22.2.0 in /usr/local/lib/python3.10/dist-packages (from jsonschema>=2.6->nbformat->otter-grader) (24.3.0)

Requirement already satisfied: jsonschema-specifications<=2023.03.6 in /usr/local/lib/python3.10/dist-packages (from jsonschema<=2.6->nbformat->otter-grader) (2024.10.1)

Requirement already satisfied: referencing<=0.28.4 in /usr/local/lib/python3.10/dist-packages (from jsonschema<=2.6->nbformat->otter-grader) (0.35.1)

Requirement already satisfied: rpds-py<=0.7.1 in /usr/local/lib/python3.10/dist-packages (from jsonschema<=2.6->nbformat->otter-grader) (0.22.3)

Requirement already satisfied: platformdirs<=2.5 in /usr/local/lib/python3.10/dist-packages (from jupyter-core<=4.7->nbconvert->otter-grader) (4.3.6)

Requirement already satisfied: jupyter-client<=6.1.12 in /usr/local/lib/python3.10/dist-packages (from nbclient<=0.5.0->nbconvert->otter-grader) (6.1.12)

Requirement already satisfied: ptyprocess<=0.5 in /usr/local/lib/python3.10/dist-packages (from pexpect<4.3->ipython->otter-grader) (0.7.0)

Requirement already satisfied: wcwidth in /usr/local/lib/python3.10/dist-packages (from prompt-toolkit!=3.0.0,!<3.0.1,<3.1.0,>=2.0.0->ipython->otter-grader) (0.2.13)

Requirement already satisfied: six<=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil<=2.8.2->pandas->otter-grader) (1.17.0)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests<=2.26.0->docker->otter-grader) (3.4.0)

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests<=2.26.0->docker->otter-grader) (3.10)

Requirement already satisfied: certifi<=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests<=2.26.0->docker->otter-grader) (2024.12.14)

Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.10/dist-packages (from beautifulsoup4->nbconvert->otter-grader) (2.6)

Requirement already satisfied: pyzmq<=13 in /usr/local/lib/python3.10/dist-packages (from jupyter-client<=6.1.12->nbclient<=0.5.0->nbconvert->otter-grader) (24.0.1)

Requirement already satisfied: wget in /usr/local/lib/python3.10/dist-packages (3.2)

Command failed with exit code 256

Out[66]: 256

```
In [67]: # Initialize Otter
import otter
grader = otter.Notebook()
```

```
%%\newcommand{\vect}[1]{\mathbf{#1}} \newcommand{\cnt}[1]{\sharp(#1)} \newcommand{\argmax}[1]{\underset{#1}{\operatorname{argmax}}}} \newcommand{\softmax}{\operatorname{softmax}} \newcommand{\Prob}{\Pr} \newcommand{\given}{\,|\,}
```

Course 236299

Project segment 1: Text classification

In this project segment you will build several varieties of text classifiers using `PyTorch`.

1. A majority baseline.
2. A naive Bayes classifier.
3. A logistic regression (single-layer perceptron) classifier.
4. A multilayer perceptron classifier.

You'll use these to classify air travel queries from a standard dataset, and compare the results.

Preparation {-}

```
In [68]: import copy
import re
import wget
import csv
import torch
import torch.nn as nn
import datasets

from datasets import load_dataset
from tokenizers import Tokenizer
from tokenizers.pre_tokenizers import Whitespace
from tokenizers import normalizers
from tokenizers.models import WordLevel
from tokenizers.trainers import WordLevelTrainer
from transformers import PreTrainedTokenizerFast
from collections import Counter
from torch import optim
from tqdm.auto import tqdm
```

```
In [69]: # Random seed
random_seed = 1234
torch.manual_seed(random_seed)

## GPU check
device = torch.device("cuda" if torch.cuda.is_available() else
                      "cpu")
print(device)
```

cuda

The task: Answer types for ATIS queries

For this and future project segments, you will be working with a standard natural-language-processing dataset, the [ATIS \(Airline Travel Information System\) dataset](#). This dataset is composed of queries about flights – their dates, times, locations, airlines, and the like.

The ATIS dataset was generated using a ["Wizard of Oz" methodology](#), a common approach in early NLP research and human-computer interaction more generally. [Subjects were asked](#) to interact with a "prototype of a voice-input information retrieval system. It has the same information that is contained in the Official Airline Guide (OAG) to help you make air travel plans." In reality, behind the curtain, two confederates were transcribing the queries and providing the answers. In this way, an "ecologically realistic" set of queries was obtained.

Over the years, the dataset has been annotated in all kinds of ways, with parts of speech, informational chunks, parse trees, and even corresponding SQL database queries. You'll use various of these annotations in future assignments. For this project segment, however, you'll pursue an easier classification task: **given a query, predict the answer type**.

These queries ask for different types of answers, such as

- Flight IDs: "Show me the flights from Washington to Boston"
- Fares: "How much is the cheapest flight to Milwaukee"
- City names: "Where does flight 100 fly to?"

In all, there are some 30 answer types to the queries.

Below is an example taken from this dataset:

Query:

```
show me the afternoon flights from washington to boston
```

SQL:

```
SELECT DISTINCT flight_1.flight_id FROM flight flight_1 ,
airport_service airport_service_1 , city city_1 ,
airport_service airport_service_2 , city city_2
  WHERE flight_1.departure_time BETWEEN 1200 AND 1800
        AND ( flight_1.from_airport =
airport_service_1.airport_code
            AND airport_service_1.city_code = city_1.city_code
            AND city_1.city_name = 'WASHINGTON'
            AND flight_1.to_airport =
airport_service_2.airport_code
            AND airport_service_2.city_code = city_2.city_code
            AND city_2.city_name = 'BOSTON' )
```

In this project segment, we will consider the answer type for a natural-language query to be the target field of the corresponding SQL query. For the above example, the answer type would be *flight_id*.

Loading and preprocessing the data

Read over this section, executing the cells, and **making sure you understand what's going on before proceeding to the next parts**.

First, let's download the dataset.

```
In [70]: data_dir = "https://raw.githubusercontent.com/nlp-236299/data/master/ATIS/"
os.makedirs('data', exist_ok=True)
for split in ['train', 'dev', 'test']:
    wget.download(f"{data_dir}/{split}.nl", out='data/')
    wget.download(f"{data_dir}/{split}.sql", out='data/')
```

Next, we process the dataset by extracting answer types from SQL queries and saving in CSV format.

```
In [71]: def get_label_from_query(query):
    """Returns the answer type from `query` by dead reckoning.
    It's basically the second or third token in the SQL query.
    """
    match = re.match(r'\s*SELECT\s+(DISTINCT\s*)?(\w+\.)?(?P<label>\w+)', query)
    if match:
        label = match.group('label')
    else:
        raise RuntimeError(f'no label in query {query}')
    return label

    for split in ['train', 'dev', 'test']:
        sql_file = f'data/{split}.sql'
        nl_file = f'data/{split}.nl'
        out_file = f'data/{split}.csv'

        with open(nl_file) as f_nl:
            with open(sql_file) as f_sql:
                with open(out_file, 'w') as fout:
                    writer = csv.writer(fout)
                    writer.writerow(('label', 'text'))
                    for text, sql in zip(f_nl, f_sql):
                        text = text.strip()
                        sql = sql.strip()
                        label = get_label_from_query(sql)
                        writer.writerow((label, text))
```

Let's take a look at what each data file looks like.

```
In [72]: shell('head "data/train.csv"')
```

```
label,text
flight_id,list all the flights that arrive at general mitchell international from
various cities
flight_id,give me the flights leaving denver august ninth coming back to boston
flight_id,what flights from tacoma to orlando on saturday
fare_id,what is the most expensive one way fare from boston to atlanta on america
n airlines
flight_id,what flights return from denver to philadelphia on a saturday
flight_id,can you list all flights from chicago to milwaukee
flight_id,show me the flights from denver that go to pittsburgh and then atlanta
flight_id,i'd like to see flights from baltimore to atlanta that arrive before noon
on and i'd like to see flights from denver to atlanta that arrive before noon
flight_id,do you have an 819 flight from denver to san francisco
```

It's a CSV file with the answer type in the first column and the query text in the second.

We use `datasets` to prepare the data, as in lab 1-5. More information on `datasets` can be found at <https://huggingface.co/docs/datasets/loading>.

```
In [73]: atis = load_dataset('csv', data_files={'train': 'data/train.csv', \
        'val': 'data/dev.csv', \
        'test': 'data/test.csv'})
```

Generating train split: 0 examples [00:00, ? examples/s]

Generating val split: 0 examples [00:00, ? examples/s]
 Generating test split: 0 examples [00:00, ? examples/s]

In [74]: `atis`

```
Out[74]: DatasetDict({
  train: Dataset({
    features: ['label', 'text'],
    num_rows: 4379
  })
  val: Dataset({
    features: ['label', 'text'],
    num_rows: 491
  })
  test: Dataset({
    features: ['label', 'text'],
    num_rows: 448
  })
})
```

```
In [75]: train_data = atis['train']
val_data = atis['val']
test_data = atis['test']

train_data.shuffle(seed=random_seed)
```

```
Out[75]: Dataset({
  features: ['label', 'text'],
  num_rows: 4379
})
```

We build a tokenizer from the training data to tokenize text and convert tokens into word ids.

```
In [76]: MIN_FREQ = 3 # words appearing fewer than 3 times are treated as 'unknown'
unk_token = '[UNK]'
pad_token = '[PAD]'

tokenizer = Tokenizer(WordLevel(unk_token=unk_token))
tokenizer.pre_tokenizer = Whitespace()
tokenizer.normalizer = normalizers.Lowercase()

trainer = WordLevelTrainer(min_frequency=MIN_FREQ, special_tokens=[pad_token, unk_token])
tokenizer.train_from_iterator(train_data['text'], trainer=trainer)
```

We use `datasets.Dataset.map` to convert text into word ids. As shown in lab 1-5, first we need to wrap `tokenizer` with the `transformers.PreTrainedTokenizerFast` class to be compatible with the `datasets` library.

```
In [77]: hf_tokenizer = PreTrainedTokenizerFast(tokenizer_object=tokenizer,
                                                pad_token=pad_token,
                                                unk_token=unk_token)
```

```
In [78]: def encode(example):
  return hf_tokenizer(example['text'])

train_data = train_data.map(encode)
```

```
val_data = val_data.map(encode)
test_data = test_data.map(encode)
```

```
Map: 0%|          | 0/4379 [00:00<?, ? examples/s]
Map: 0%|          | 0/491 [00:00<?, ? examples/s]
Map: 0%|          | 0/448 [00:00<?, ? examples/s]
```

We also need to convert label strings into label ids.

```
In [79]: # Add a new column `label_id`
train_data = train_data.add_column('label_id', train_data['label'])
val_data = val_data.add_column('label_id', val_data['label'])
test_data = test_data.add_column('label_id', test_data['label'])

# Convert feature `label_id` from strings to integer ids
train_data = train_data.class_encode_column('label_id')

# Use the label vocabulary on training data to convert val and test sets
label2id = train_data.features['label_id']._str2int
val_data = val_data.class_encode_column('label_id')
val_data = val_data.align_labels_with_mapping(label2id, "label_id")
test_data = test_data.class_encode_column('label_id')
test_data = test_data.align_labels_with_mapping(label2id, "label_id")
```

```
Casting to class labels: 0%|          | 0/4379 [00:00<?, ? examples/s]
Casting to class labels: 0%|          | 0/491 [00:00<?, ? examples/s]
Aligning the labels: 0%|          | 0/491 [00:00<?, ? examples/s]
Casting to class labels: 0%|          | 0/448 [00:00<?, ? examples/s]
Aligning the labels: 0%|          | 0/448 [00:00<?, ? examples/s]
```

```
In [80]: # Compute size of vocabulary
text_vocab = tokenizer.get_vocab()
label_vocab = train_data.features['label_id']._str2int
vocab_size = len(text_vocab)
num_labels = len(label_vocab)
print(f"Size of vocab: {vocab_size}")
print(f"Number of labels: {num_labels}")
```

```
Size of vocab: 514
Number of labels: 30
```

To get a sense of the kinds of things that are asked about in this dataset, here is the list of all of the answer types in the training data, along with their label IDs.

```
In [81]: for label in label_vocab:
print(f"{label_vocab[label]:2d} {label}")
```

```

0 advance_purchase
1 aircraft_code
2 airline_code
3 airport_code
4 airport_location
5 arrival_time
6 basic_type
7 booking_class
8 city_code
9 city_name
10 count
11 day_name
12 departure_time
13 fare_basis_code
14 fare_id
15 flight_id
16 flight_number
17 ground_fare
18 meal_code
19 meal_description
20 miles_distant
21 minimum_connect_time
22 minutes_distant
23 restriction_code
24 state_code
25 stop_airport
26 stops
27 time_elapsed
28 time_zone_code
29 transport_type

```

Handling unknown words

Note that we mapped words appearing fewer than 3 times to a special *unknown* token (we're using `[UNK]`). We do this for two reasons:

1. Due to the scarcity of such rare words in training data, we might not be able to learn generalizable conclusions about them.
2. Introducing an unknown token allows us to deal with out-of-vocabulary words in the test data as well: we just map those words to `[UNK]`.

```

In [82]: print (f"Unknown token: {unk_token}")
unk_index = text_vocab[unk_token]
print (f"Unknown token id: {unk_index}")

# UNK example
example_unk_token = 'IAmAnUnknownWordForSure'
print (f"An unknown token: {example_unk_token}")
print (f"Mapped back to word id: {hf_tokenizer(example_unk_token).input_ids}")
print (f"Mapped to [UNK]'s?: {all([id == unk_index for id in hf_tokenizer(exempl

```

```

Unknown token: [UNK]
Unknown token id: 1
An unknown token: IAmAnUnknownWordForSure
Mapped back to word id: [1]
Mapped to [UNK]'s?: True

```

To facilitate batching sentences of different lengths into the same tensor as we'll see later, we also reserved a special padding symbol `[PAD]`.

```
In [83]: print (f"Padding token: {pad_token}")
         pad_index = text_vocab[pad_token]
         print (f"Padding token id: {pad_index}")
```

Padding token: [PAD]

Padding token id: 0

Batching the data

To load data in batches, we use `torch.utils.data.DataLoader`. This enables us to iterate over the dataset under a given `BATCH_SIZE` which specifies how many examples we want to process at a time.

```
In [84]: BATCH_SIZE = 32

# Defines how to batch a list of examples together
def collate_fn(examples):
    batch = {}
    bsz = len(examples)
    label_ids = []
    for example in examples:
        label_ids.append(example['label_id'])
    label_batch = torch.LongTensor(label_ids).to(device)
    input_ids = []
    for example in examples:
        input_ids.append(example['input_ids'])
    max_length = max([len(word_ids) for word_ids in input_ids])
    text_batch = torch.zeros(bsz, max_length).long().fill_(pad_index).to(device)
    for b in range(bsz):
        text_batch[b][:len(input_ids[b])] = torch.LongTensor(input_ids[b]).to(device)

    batch['label_ids'] = label_batch
    batch['input_ids'] = text_batch
    return batch

train_iter = torch.utils.data.DataLoader(train_data, batch_size=BATCH_SIZE, coll
val_iter = torch.utils.data.DataLoader(val_data, batch_size=BATCH_SIZE, coll
test_iter = torch.utils.data.DataLoader(test_data, batch_size=BATCH_SIZE, coll
```

Let's look at a single batch from one of these iterators.

```
In [85]: batch = next(iter(train_iter))
         text = batch['input_ids']
         print (f"Size of text batch: {text.size()}")
         print (f"Third sentence in batch: {text[2]}")
         print (f"Mapped back to string: {hf_tokenizer.decode(text[2])}")
         print (f"Mapped back to string skipping padding: {hf_tokenizer.decode(text[2], s

         label = batch['label_ids']
         label_vocab_itos = train_data.features['label_id']._int2str # map from label ids
         print (f"Size of label batch: {label.size()}")
         print (f"Third label in batch: {label[2]}")
         print (f"Mapped back to string: {label_vocab_itos[label[2].item()]})")
```

```
Size of text batch: torch.Size([32, 31])
Third sentence in batch: tensor([ 7,  4,  3, 180,  2, 114,  6, 119,  0,
  0,  0,  0,  0,  0,
        0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0], device='cuda:0')
Mapped back to string: what flights from tacoma to orlando on saturday [PAD] [PA
D] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD]
[PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD] [PAD]
Mapped back to string skipping padding: what flights from tacoma to orlando on sa
turday
Size of label batch: torch.Size([32])
Third label in batch: 15
Mapped back to string: flight_id
```


- $\mathbf{w} = w_1 \cdots w_M$: A text to be classified, each element w_j being a word token.
- $\mathbf{v} = \{v_1, \dots, v_V\}$: A vocabulary, each element v_k being a word type.
- $\mathbf{x} = \langle x_1, \dots, x_X \rangle$: Input features to a model.
- $\mathbf{y} = \{y_1, \dots, y_N\}$: The output classes of a model, each element y_i being a class label.
- $\mathbf{T} = \langle \mathbf{w}^{(1)}, \dots, \mathbf{w}^{(T)} \rangle$: The training corpus of texts.
- $\mathbf{Y} = \langle \mathbf{y}^{(1)}, \dots, \mathbf{y}^{(T)} \rangle$: The corresponding gold labels for the training examples in T .

To Do: Establish a majority baseline

A simple baseline for classification tasks is to always predict the most common class. Given a training set of texts \mathbf{T} labeled by classes \mathbf{Y} , we classify an input text $\mathbf{w} = w_1 \cdots w_M$ as the class y_i that occurs most frequently in the training data, that is, specified by

$$\underset{i}{\operatorname{argmax}} \#(y_i)$$

and thus ignoring the input \mathbf{x} entirely (!).

Implement the majority baseline and compute test accuracy using the starter code below. For this baseline, and for the naive Bayes classifier later, we don't need to use the validation set since we don't tune any hyper-parameters.

```
In [87]: def majority_baseline_accuracy(train_data, test_data):
    label_counts = Counter(train_data['label'])
    most_common_label = label_counts.most_common(1)[0][0]

    correct = sum(1 for label in test_data['label'] if label == most_common_label)
    accuracy = correct / len(test_data['label'])

    return most_common_label, accuracy
```

How well does your classifier work? Let's see:

```
In [88]: # Call the method to establish a baseline
most_common_label, test_accuracy = majority_baseline_accuracy(train_data, test_data)
# For comparison, evaluate it on the training data as well
_, train_accuracy = majority_baseline_accuracy(train_data, train_data)

print(f'Most common label: {most_common_label}\n'
      f'Test accuracy:      {test_accuracy:.3f}\n'
      f'Train accuracy:     {train_accuracy:.3f}')
```

```
Most common label: flight_id
Test accuracy:      0.683
Train accuracy:     0.733
```

To Do: Implement a Naive Bayes classifier

Review of the naive Bayes method

Recall from lab 1-3 that the Naive Bayes classification method classifies a text $\mathbf{w} = \langle w_1, w_2, \dots, w_M \rangle$ as the class y_i given by the following maximization:

$$\operatorname{argmax}_i \Pr(y_i | \mathbf{w}) \approx \operatorname{argmax}_i \Pr(y_i) \cdot \prod_{j=1}^M \Pr(w_j | y_i)$$

or equivalently (since taking the log is monotonic)

$$\operatorname{argmax}_i \Pr(y_i | \mathbf{w}) = \operatorname{argmax}_i \log \Pr(y_i | \mathbf{w}) \quad (1)$$

$$\approx \operatorname{argmax}_i \left(\log \Pr(y_i) + \sum_{j=1}^M \log \Pr(w_j | y_i) \right) \quad (2)$$

All we need, then, to apply the Naive Bayes classification method is values for the various log probabilities: the priors $\log \Pr(y_i)$ and the likelihoods $\log \Pr(w_j | y_i)$, for each feature (word) w_j and each class y_i .

We can estimate the prior probabilities $\Pr(y_i)$ by examining the empirical probability in the training set. That is, we estimate

$$\Pr(y_i) \approx \frac{\#(y_i)}{\sum_j \#(y_j)}$$

We can estimate the likelihood probabilities $\Pr(w_j | y_i)$ similarly by examining the empirical probability in the training set. That is, we estimate

$$\Pr(w_j | y_i) \approx \frac{\#(w_j, y_i)}{\sum_{j'} \#(w_{j'}, y_i)}$$

To allow for cases in which the count $\#(w_j, y_i)$ is zero, we can use a modified estimate incorporating add- δ smoothing:

$$\Pr(w_j | y_i) \approx \frac{\#(w_j, y_i) + \delta}{\sum_{j'} \#(w_{j'}, y_i) + \delta \cdot V}$$

Two conceptions of the naive Bayes method implementation

We can store all of these parameters in different ways, leading to two different implementation conceptions. We review two conceptions of implementing the naive Bayes classification of a text $\mathbf{w} = \langle w_1, w_2, \dots, w_M \rangle$, corresponding to using different representations of the input \mathbf{x} to the model: the index representation and the bag-of-words representation.

Within each conception, the parameters of the model will be stored in one or more matrices. The conception dictates what operations will be performed with these matrices.

Using the index representation

In the first conception, we take the input elements $\mathbf{x} = \langle x_1, x_2, \dots, x_M \rangle$ to be the *vocabulary indices* of the words $\mathbf{w} = w_1 \cdots w_M$. That is, each word token w_i is of the word type in the vocabulary \mathbf{v} at index x_i , so

$$v_{x_i} = w_i$$

In this representation, the input vector has the same length as the word sequence.

We think of the likelihood probabilities as forming a matrix, call it \mathbf{L} , where the i, j -th element stores $\log \Pr(v_j | y_i)$.

$$\mathbf{L}_{ij} = \log \Pr(v_j | y_i)$$

Similarly, for the priors, we'll have

$$\mathbf{P}_i = \log \Pr(y_i)$$

Now the maximization can be implemented as

$$\operatorname{argmax}_i \log \Pr(y_i) + \sum_{j=1}^M \log \Pr(w_j | y_i) = \operatorname{argmax}_i \mathbf{P}_i + \sum_{j=1}^M \mathbf{L}_{i,x_j} \quad (3)$$

Implemented in this way, we see that the use of each input x_i is as an *index* into the likelihood matrix.

Using the bag-of-words representation

Notice that since each word in the input is treated separately, the order of the words doesn't matter. Rather, all that matters is how frequently each word type occurs in a text. Consequently, we can use the bag-of-words representation introduced in lab 1-1.

Recall that the bag-of-words representation of a text is just its frequency distribution over the vocabulary, which we will notate $\text{bow}(\mathbf{w})$. Given a vocabulary of word types $\mathbf{v} = \langle v_1, v_2, \dots, v_V \rangle$, the representation of a sentence $\mathbf{w} = \langle w_1, w_2, \dots, w_M \rangle$ is a vector \mathbf{x} of size V , where

$$\text{bow}(\mathbf{w})_j = \sum_{i=1}^M 1[w_i = v_j] \quad \text{for } 1 \leq j \leq V$$

We write $1[w_i = v_j]$ to indicate 1 if $w_i = v_j$ and 0 otherwise. For convenience, we'll add an extra $(V + 1)$ -st element to the end of the bag-of-words vector, a single 1 whose use will be clear shortly. That is,

$$\text{bow}(\mathbf{w})_{V+1} = 1$$

Under this conception, then, we'll take the input \mathbf{x} to be $\text{bow}(\mathbf{w})$. Instead of the input having the same length as the text, it has the same length as the vocabulary.

As described in lecture, represented in this way, the quantity to be maximized in the naive Bayes method

$$\log \Pr(y_i) + \sum_{j=1}^M \log \Pr(w_j | y_i)$$

can be calculated as $\hat{y} = \underset{i}{\operatorname{argmax}} \tilde{y}_i$:1


$$\log \Pr(y_i) + \sum_{j=1}^V x_j \cdot \log \Pr(v_j | y_i)$$

which is just $\mathbf{U}\mathbf{x}$ for a suitable choice of $N \times (V + 1)$ matrix \mathbf{U} , namely

$$\mathbf{U}_{ij} = \begin{cases} \log \Pr(v_j | y_i) & 1 \leq i \leq N \text{ and } 1 \leq j \leq V \\ \log \Pr(y_i) & 1 \leq i \leq N \text{ and } j = V + 1 \end{cases}$$

Under this implementation conception, we've reduced naive Bayes calculations to a single matrix operation. This conception is depicted in the figure at right.

You are free to use either conception in your implementation of naive Bayes.



Implement the naive Bayes classifier

For the implementation, we ask you to implement a Python class `NaiveBayes` that will have (at least) the following three methods:

1. `__init__` : An initializer that takes `text_vocab`, `label_vocab`, and `pad_index` as inputs.
2. `train` : A method that takes a training data iterator and estimates all of the log probabilities $\log \Pr(y_i)$ and $\log \Pr(v_j | y_i)$ as described above. Perform add- δ

smoothing with $\delta = 1$. These parameters will be used by the `evaluate` method to evaluate a test dataset for accuracy, so you'll want to store them in some data structures in objects of the class.

3. `evaluate` : A method that takes a test data iterator and evaluates the accuracy of the trained model on the test set.

You can organize your code using either of the conceptions of Naive Bayes described above.

You should expect to achieve about an **86% test accuracy** on the ATIS task.

```
In [92]: class NaiveBayes():
    def __init__(self, text_vocab, label_vocab, pad_index):
        self.pad_index = pad_index
        self.V = len(text_vocab) # vocabulary size
        self.N = len(label_vocab) # the number of classes
        self.priors = None
        self.likelihoods = None

    def train(self, iterator):
        """Calculates and stores log probabilities for training dataset `iterator`"""
        label_counts = Counter()
        word_label_counts = Counter()
        total_labels = 0

        for batch in iterator:
            labels = batch['label_ids']
            inputs = batch['input_ids']

            for i, label in enumerate(labels):
                label_counts[label.item()] += 1
                total_labels += 1
                for token in inputs[i]:
                    if token.item() != self.pad_index:
                        word_label_counts[(token.item(), label.item())] += 1

        self.priors = {
            label: count / total_labels for label, count in label_counts.items()
        }

        self.likelihoods = {}
        for (word, label), count in word_label_counts.items():
            self.likelihoods[(word, label)] = (count + 1) / (
                label_counts[label] + self.V
            )

    def evaluate(self, iterator):
        """Returns the model's accuracy on a given dataset `iterator`."""
        correct = 0
        total = 0

        for batch in iterator:
            labels = batch['label_ids']
            inputs = batch['input_ids']

            for i, label in enumerate(labels):
                scores = {
```

```

        y: torch.log(torch.tensor(self.priors.get(y, 1e-9)))
    for y in range(self.N)
}

for token in inputs[i]:
    if token.item() != self.pad_index:
        for y in range(self.N):
            scores[y] += torch.log(
                torch.tensor(self.likelihoods.get((token.item(),
                    )

predicted_label = max(scores, key=scores.get)
if predicted_label == label.item():
    correct += 1
total += 1

return correct / total

```

```

In [93]: # Instantiate and train classifier
nb_classifier = NaiveBayes(text_vocab, label_vocab, pad_index)
nb_classifier.train(train_iter)

# Evaluate model performance
print(f'Training accuracy: {nb_classifier.evaluate(train_iter):.3f}\n'
      f'Test accuracy:      {nb_classifier.evaluate(test_iter):.3f}')

```

```

Training accuracy: 0.817
Test accuracy:      0.777

```

To Do: Implement a logistic regression classifier

In this part, you'll complete a PyTorch implementation of a logistic regression (equivalently, a single layer perceptron) classifier. We review logistic regression here highlighting the similarities to the matrix-multiplication conception of naive Bayes. Thus, we take the input \mathbf{x} to be the bag-of-words representation $\text{bow}(\mathbf{w})$. But as before you are free to use either implementation approach.

Review of logistic regression

Similar to naive Bayes, in logistic regression, we assign a probability to a text \mathbf{x} by merely multiplying an $N \times V$ matrix \mathbf{U} by it. However, we don't stipulate that the values in the matrix \mathbf{U} be estimated from the training corpus in the "naive Bayes" manner. Instead, we allow them to take on any value, using a training regime to select good values.

In order to make sure that the output of the matrix multiplication $\mathbf{U}\mathbf{x}$ is mapped onto a probability distribution, we apply a nonlinear function to renormalize the values. We use the softmax function, a generalization of the sigmoid function from lab 1-4, defined by

$$\text{softmax}(\mathbf{z})_i = \frac{\exp(z_i)}{\sum_{j=1}^N \exp(z_j)}$$

for each of the indices i from 1 to N .

In summary, we model $\Pr(y | \mathbf{x})$ as


$$\Pr(y_i | \mathbf{x}) = \text{softmax}(\mathbf{U}\mathbf{x})_i$$

The calculation of $\Pr(y | \mathbf{x})$ for each text \mathbf{x} is referred to as the *forward* computation. In summary, the forward computation for logistic regression involves a linear calculation ($\mathbf{U}\mathbf{x}$) followed by a nonlinear calculation (softmax). We think of the perceptron (and more generally many of these neural network models) as transforming from one representation to another. A perceptron performs a linear transformation from the index or bag-of-words representation of the text to a representation as a vector, followed by a nonlinear transformation, a softmax or sigmoid, giving a representation as a probability distribution over the class labels. This single-layer perceptron thus

involves two *sublayers*. (In the next part of the project segment, you'll experiment with a multilayer perceptron, with two perceptron layers, and hence four sublayers.)

The loss function you'll use is the negative log probability $-\log \Pr(y | \mathbf{x})$. The negative is used, since it is convention to minimize loss, whereas we want to maximize log likelihood.

The forward and loss computations are illustrated in the figure at right. In practice, for numerical stability reasons, PyTorch absorbs the softmax operation into the loss function `nn.CrossEntropyLoss`. That is, the input to the `nn.CrossEntropyLoss` function is



the vector of sums $\mathbf{U}\mathbf{x}$ (the last step in the box marked "your job" in the figure) rather than the vector of probabilities $\Pr(y | \mathbf{x})$. That makes things easier for you (!), since you're responsible only for the first sublayer.

Given a forward computation, the weights can then be adjusted by taking a step opposite to the gradient of the loss function. Adjusting the weights in this way is referred to as the *backward* computation. Fortunately, `torch` takes care of the backward computation for you, just as in lab 1-5.

The optimization process of performing the forward computation, calculating the loss, and performing the backward computation to improve the weights is done repeatedly until the process converges on a (hopefully) good set of weights. You'll find this optimization process in the `train_all` method that we've provided. The trained weights can then be used to perform classification on a test set. See the `evaluate` method.

Implement the logistic regression classifier

For the implementation, we ask you to implement a logistic regression classifier as a subclass of `torch.nn.Module`. You need to implement the following methods:

1. `__init__`: an initializer that takes `text_vocab`, `label_vocab`, and `pad_index` as inputs.

During initialization, you'll want to define a `tensor` of weights, wrapped in `torch.nn.Parameter`, *initialized randomly*, which plays the role of \mathbf{U} . The elements of this tensor are the parameters of the `torch.nn` instance in the following special technical sense: It is the parameters of the module whose gradients will be calculated and whose values will be updated. Alternatively, **you might find it easier** to use the `nn.Embedding` *module* which is a wrapper to the weight tensor with a lookup implementation.

2. `forward`: given a text batch of size `batch_size X max_length`, return a tensor of logits of size `batch_size X num_labels`. That is, for each text \mathbf{x} in the batch and each label y , you'll be calculating $\mathbf{U}\mathbf{x}$ as shown in the figure, returning a tensor of these values. Note that the softmax operation is absorbed into `nn.CrossEntropyLoss` so you won't need to deal with that.
3. `train_all`: A method that performs training. You might find lab 1-5 useful.
4. `evaluate`: A method that takes a test data iterator and evaluates the accuracy of the trained model on the test set.

Some things to consider:

1. The parameters of the model, the weights, need to be initialized properly. We suggest initializing them to some small random values. See `torch.uniform_`.

2. You'll want to make sure that padding tokens are handled properly. What should the weight be for the padding token?
3. In extracting the proper weights to sum up, based on the word types in a sentence, we are essentially doing a lookup operation. You might find `nn.Embedding` or `torch.gather` useful.

You should expect to achieve about **90%** accuracy on the ATIS classification task.

```
In [94]: class LogisticRegression(nn.Module):
def __init__(self, text_vocab, label_vocab, pad_index):
    super().__init__()
    self.pad_index = pad_index
    # Keep the vocabulary sizes available
    self.N = len(label_vocab) # num_classes
    self.V = len(text_vocab) # vocab_size
    # Specify cross-entropy loss for optimization
    self.criterion = nn.CrossEntropyLoss()
    # TODO: Create and initialize a tensor for the weights,
    #         or create an nn.Embedding module and initialize
    #         Matrice de poids (Ux), initialisée aléatoirement
    self.weights = nn.Embedding(self.V, self.N)
    nn.init.uniform_(self.weights.weight, -0.1, 0.1)

def forward(self, text_batch):
    # TODO: Calculate the Logits (Ux) for the `text_batch`,
    #         returning a tensor of size batch_size x num_labels
    #         Obtenir les poids pour chaque mot dans le batch
    embeddings = self.weights(text_batch) # [batch_size x seq_len x num_labels]

    # Somme sur la séquence pour obtenir [batch_size x num_labels]
    logits = embeddings.sum(dim=1) # Somme des vecteurs de chaque mot
    return logits

def train_all(self, train_iter, val_iter, epochs=8, learning_rate=3e-3):
    # Use Adam to optimize the parameters
    optim = torch.optim.Adam(self.parameters(), lr=learning_rate)
    best_validation_accuracy = -float('inf')
    best_model = None

    # Run the optimization for multiple epochs
    with tqdm(range(epochs), desc='train', position=0) as pbar:
        for epoch in pbar:

            # Switch the module to training mode (each epoch, since
            # `self.evaluate` call below resets it to evaluation mode)
            self.train()

            c_num = 0
            total = 0
            running_loss = 0.0

            for batch in tqdm(train_iter, desc='batch', leave=False):
                # TODO: set labels, compute logits (Ux in this model),
                #         loss, and update parameters
                labels = batch['label_ids'] # Les labels du batch
                inputs = batch['input_ids'] # Les phrases tokenisées

                # Réinitialisation des gradients
```

```

optim.zero_grad()

        # Calcul des logits et de la perte
logits = self.forward(inputs) # [batch_size x num_labels]
loss = self.criterion(logits, labels)
loss.backward()
optim.step()
# Prepare to compute the accuracy
predictions = torch.argmax(logits, dim=1)
total += predictions.size(0)
c_num += (predictions == labels).float().sum().item()
running_loss += loss.item() * predictions.size(0)

# Evaluate and track improvements on the validation dataset
validation_accuracy = self.evaluate(val_iter)
if validation_accuracy > best_validation_accuracy:
    best_validation_accuracy = validation_accuracy
    self.best_model = copy.deepcopy(self.state_dict())
epoch_loss = running_loss / total
epoch_acc = c_num / total
pbar.set_postfix(epoch=epoch+1, loss=epoch_loss, train_acc = epoch_acc,

def evaluate(self, iterator):
    """Returns the model's accuracy on a given dataset `iterator`."""
    self.eval() # switch the module to evaluation mode
    # TODO: Compute accuracy
    self.eval() # Met le modèle en mode évaluation
    correct = 0
    total = 0

    with torch.no_grad():
        for batch in iterator:
            labels = batch['label_ids']
            inputs = batch['input_ids']

            logits = self.forward(inputs)
            predictions = torch.argmax(logits, dim=1)
            total += predictions.size(0)
            correct += (predictions == labels).float().sum().item()

    return correct / total

```

```

In [95]: # Instantiate the logistic regression classifier and run it
model = LogisticRegression(text_vocab, label_vocab, pad_index).to(device)
model.train_all(train_iter, val_iter)
model.load_state_dict(model.best_model)
test_accuracy = model.evaluate(test_iter)
print (f'Test accuracy: {test_accuracy:.4f}')

```

```

train:  0%|          | 0/8 [00:00<?, ?it/s]
batch:  0%|          | 0/137 [00:00<?, ?it/s]
batch:  0%|          | 0/137 [00:00<?, ?it/s]
batch:  0%|          | 0/137 [00:00<?, ?it/s]
batch:  0%|          | 0/137 [00:00<?, ?it/s]
batch:  0%|          | 0/137 [00:00<?, ?it/s]
batch:  0%|          | 0/137 [00:00<?, ?it/s]
batch:  0%|          | 0/137 [00:00<?, ?it/s]
batch:  0%|          | 0/137 [00:00<?, ?it/s]
Test accuracy: 0.9219

```


To Do: Implement a multilayer perceptron

Review of multilayer perceptrons

In the last part, you implemented a perceptron, a model that involved a linear calculation (the sum of weights) followed by a nonlinear calculation (the softmax, which converts the summed weight values to probabilities). In a multi-layer perceptron, we take the output of the first perceptron to be the input of a second perceptron (and of course, we could continue on with a third or even more).

In this part, you'll implement the forward calculation of a two-layer perceptron, again letting PyTorch handle the backward calculation as well as the optimization of parameters. The first layer will involve a linear summation as before and a **sigmoid** as the nonlinear function. The second will involve a linear summation and a softmax (the latter absorbed, as before, into the loss function).

Thus, the difference from the logistic regression implementation is simply the adding of



the sigmoid and second linear calculations. See the figure for the structure of the computation.

Implement a multilayer perceptron classifier

For the implementation, we ask you to implement a two layer perceptron classifier, again as a subclass of the `torch.nn module`. You might reuse quite a lot of the code from logistic regression. As before, you need to implement the following methods:

1. `__init__` : An initializer that takes `text_vocab` , `label_vocab` , `pad_index` , and `hidden_size` specifying the size of the hidden layer (e.g., in the above illustration, `hidden_size` is `D`).

During initialization, you'll want to define two tensors of weights, which serve as the parameters of this model, one for each layer. You'll want to [initialize them randomly](#).

The weights in the first layer are a kind of lookup (as in the previous part), mapping words to a vector of size `hidden_size` . The `nn.Embedding module` is a good way to set up and make use of this weight tensor.

The weights in the second layer define a linear mapping from vectors of size `hidden_size` to vectors of size `num_labels` . The `nn.Linear module` or `torch.mm` for matrix multiplication may be helpful here.

2. `forward` : Given a text batch of size `batch_size X max_length` , the `forward` function returns a tensor of logits of size `batch_size X num_labels` .

That is, for each text \mathbf{x} in the batch and each label c , you'll be calculating $MLP(bow(\mathbf{x}))$ as shown in the illustration above, returning a tensor of these values. Note that the softmax operation is absorbed into `nn.CrossEntropyLoss` so you don't need to worry about that.

For the sigmoid sublayer, you might find `nn.Sigmoid` useful.

3. `train_all` : A method that performs training. You might find lab 1-5 useful.
4. `evaluate` : A method that takes a test data iterator and evaluates the accuracy of the trained model on the test set.

You should expect to achieve at least **90%** accuracy on the ATIS classification task.

```
In [96]: class MultiLayerPerceptron(nn.Module):
def __init__(self, text_vocab, label_vocab, pad_index, hidden_size=128):
    super().__init__()
    self.pad_index = pad_index
    self.hidden_size = hidden_size
    # Keep the vocabulary sizes available
    self.N = len(label_vocab) # num_classes
    self.V = len(text_vocab) # vocab_size

    # Specify cross-entropy loss for optimization
```

```

self.criterion = nn.CrossEntropyLoss()

# TODO: Create and initialize neural modules
# Embedding layer to convert words into vectors
self.embedding = nn.Embedding(self.V, hidden_size, padding_idx=pad_index)
nn.init.uniform_(self.embedding.weight, -0.1, 0.1)

# Linear layers for classification
self.fc1 = nn.Linear(hidden_size, hidden_size)
self.fc2 = nn.Linear(hidden_size, self.N)

# Activation function
self.activation = nn.ReLU()

def forward(self, text_batch):
    # TODO: Calculate the logits for the `text_batch`,
    #         returning a tensor of size batch_size x num_labels
    # Embed the input words
    embeddings = self.embedding(text_batch) # [batch_size x seq_len x hidden_size]

    # Sum over the sequence length
    summed_embeddings = embeddings.sum(dim=1) # [batch_size x hidden_size]

    # Apply the first fully connected layer and activation
    hidden_output = self.activation(self.fc1(summed_embeddings)) # [batch_size x hidden_size]

    # Apply the second fully connected layer
    logits = self.fc2(hidden_output) # [batch_size x num_labels]

    return logits

def train_all(self, train_iter, val_iter, epochs=8, learning_rate=3e-3):
    # Use Adam to optimize the parameters
    optim = torch.optim.Adam(self.parameters(), lr=learning_rate)
    best_validation_accuracy = -float('inf')
    best_model = None
    # Run the optimization for multiple epochs
    with tqdm(range(epochs), desc='train', position=0) as pbar:
        for epoch in pbar:
            # Switch the module to training mode
            self.train()
            c_num = 0
            total = 0
            running_loss = 0.0

            for batch in tqdm(train_iter, desc='batch', leave=False):
                # TODO: set labels, compute logits (Ux in this model),
                # loss, and update parameters
                labels = batch['label_ids']
                inputs = batch['input_ids']

                # Reset gradients
                optim.zero_grad()

                # Compute logits and loss
                logits = self.forward(inputs)
                loss = self.criterion(logits, labels)
                loss.backward()

```

```

        optim.step()

        # Prepare to compute the accuracy
        predictions = torch.argmax(logits, dim=1)
        total += predictions.size(0)
        c_num += (predictions == labels).float().sum().item()
        running_loss += loss.item() * predictions.size(0)

    # Evaluate on validation data
    validation_accuracy = self.evaluate(val_iter)
    if validation_accuracy > best_validation_accuracy:
        best_validation_accuracy = validation_accuracy
        self.best_model = copy.deepcopy(self.state_dict())

    epoch_loss = running_loss / total
    epoch_acc = c_num / total
    pbar.set_postfix(epoch=epoch+1, loss=epoch_loss, train_acc=epoch

def evaluate(self, iterator):
    """Returns the model's accuracy on a given dataset `iterator`."""
    # TODO: Compute accuracy
    self.eval()
    correct = 0
    total = 0

    with torch.no_grad():
        for batch in iterator:
            labels = batch['label_ids']
            inputs = batch['input_ids']

            logits = self.forward(inputs)
            predictions = torch.argmax(logits, dim=1)
            total += predictions.size(0)
            correct += (predictions == labels).float().sum().item()

    return correct / total

```

```

In [97]: # Instantiate classifier and run it
model = MultiLayerPerceptron(text_vocab, label_vocab, pad_index, hidden_size=128)
model.train_all(train_iter, val_iter)
model.load_state_dict(model.best_model)
test_accuracy = model.evaluate(test_iter)
print (f'Test accuracy: {test_accuracy:.4f}')

```

```

train:  0%|          | 0/8 [00:00<?, ?it/s]
batch:  0%|          | 0/137 [00:00<?, ?it/s]
batch:  0%|          | 0/137 [00:00<?, ?it/s]
batch:  0%|          | 0/137 [00:00<?, ?it/s]
batch:  0%|          | 0/137 [00:00<?, ?it/s]
batch:  0%|          | 0/137 [00:00<?, ?it/s]
batch:  0%|          | 0/137 [00:00<?, ?it/s]
batch:  0%|          | 0/137 [00:00<?, ?it/s]
batch:  0%|          | 0/137 [00:00<?, ?it/s]
batch:  0%|          | 0/137 [00:00<?, ?it/s]
Test accuracy: 0.9196

```

Lessons learned

Question: Rank the methods as to their performance. What lessons do you learn from the ranking?

The best model based on our experiments was the Multilayer Perceptron. After it comes Logistic Regression and Naive Bayes comes last Multilayer Perceptron outperformed others due to its utilization of an advanced architecture in conjunction with non-linear activation functions which enables it to extract better features out of the inputs. The method is Linear Regression more basic and efficient when used for linearly separable datasets, but had limited scope in terms of more complicated relationships. Our observations show Naive Bayes as the least suitable model because it relies on oversimplified assumptions regarding the independence of words which works best for other than natural language data. Clearly we can see that advanced architectures with non-linearities and deep structures are more efficient than most of their simpler counterparts; it must be noted that this also translates into more computing power and longer training periods. In some simple cases or where only limited computing resources are available, ways like Logistic Regression or even Naive Bayes can function as expected.

Question: Take a look at some of the examples that were classified correctly and incorrectly by your best method. Do you notice anything about the incorrectly classified examples that might indicate *why* they were classified incorrectly?

From what has been outlined before, it seems that the model is deficient in context and has not been sufficiently exposed to uncommon data. In the majority of the instances, the model could not comprehend due to inadequate context. Increasing the dataset or adding more diverse alternatives can mend these issues. Uncommon or ambiguous data can be included while training the model, due to which a better association can be formed. Ambiguous wordings or even questions which are multi labeled can add to the training process, providing a more fulfilling data set. An expanded and reworked dataset can aid in enhancing accuracy while labeling the sentences

In []: ...

Prompting Modern Large Language Models (LLMs)

Question: Modern large-scale language models (such as Claude, ChatGPT, Gemini, Llama) have various capabilities, that can be shown by prompting them correctly (i.e. giving them a correct input prompt). For example, they can even be useful for solving text classification, discussed in this segment. Try to see if you can prompt an LLM (of your choosing) to solve the task of text classification on some of the examples of data seen in this segment. Write a short paragraph about your experience. Note that your not expected to devise a fool-proof prompting method, but only to qualitatively experiment with prompting.

I attempted to use a large language model like ChatGPT for the text classification problem by giving it prompts containing sample sentences and sample labels. The aim of the model was to estimate the labels during use for novel sentences. The model's performance was quite reasonable, especially if the examples provided during the prompt were lucid and varied. New task often floundered on somewhat ambiguous or rare phrases, a problem that mirrors the practice models that we tested. The experience made me appreciate how important the quality of the prompt is, as it contributes a lot to the performance of the LLM. This can be especially useful to LLM models, as being verbose on the task and providing sufficient examples increases its efficiency. What this means is that with good attention to prompt design, LLMs can be used productively for tasks such as text classification without having to be explicitly trained on the dataset.

Debrief

Question: We're interested in any thoughts you have about this project segment so that we can improve it for later years, and to inform later segments for this year. Please list any issues that arose or comments you have to improve the project segment. Useful things to comment on include the following, but you are not restricted to these:

- Was the project segment clear or unclear? Which portions?
- Were the readings appropriate background for the project segment?
- Are there additions or changes you think would make the project segment better?

The project segment was clear overall, and the step-by-step instructions made it easy to follow. The readings provided sufficient background to understand the concepts and implement the methods. The practical coding tasks were well-designed and aligned with the theoretical parts, which helped reinforce the learning. To make it even better, more examples of how to handle edge cases or ambiguous data could be added, and a brief explanation of how to debug common issues might help future students. Overall, it was a well-structured and engaging project segment.

Instructions for submission of the project segment

This project segment should be submitted to Gradescope, which will be made available some time before the due date.

Project segment notebooks are manually graded, not autograded using otter as labs are. (Otter is used within project segment notebooks to synchronize distribution and solution code however.) **We will not run your notebook before grading it.** Instead, we ask that you submit the already freshly run notebook. The best method is to "restart kernel and run all cells", allowing time for all cells to be run to completion. You should submit your code to Gradescope at the code submission assignment at https://www.gradescope.com/courses/903849?submit_assignment_id=5229513.

We also request that you **submit a PDF of the freshly run notebook**. The simplest method is to use "Export notebook to PDF", which will render the notebook to PDF via LaTeX. If that doesn't work, the method that seems to be most reliable is to export the notebook as HTML (if you are using Jupyter Notebook, you can do so using `File -> Print Preview`), open the HTML in a browser, and print it to a file. Then make sure to add the file to your git commit. Please name the file the same name as this notebook, but with a `.pdf` extension. (Conveniently, the methods just described will use that name by default.) You can then perform a git commit and push and submit the commit to Gradescope at https://www.gradescope.com/courses/903849?submit_assignment_id=5229489.

End of project segment 1 {-}